

## EXPERIMENT -1

### **Aim:**

- Load data in Pandas.
- Description of the dataset.
- Drop columns that aren't useful.
- Drop rows with maximum missing values.
- Take care of missing data.
- Create dummy variables.
- Find out outliers (manually)
- standardization and normalization of columns

### **Data preprocessing**

Data preprocessing is the process of transforming raw data into an understandable format. It is also an important step in data mining as we cannot work with raw data. The quality of the data should be checked before applying machine learning or data mining algorithms.

### **Why is Data Preprocessing important?**

Preprocessing of data is mainly to check the data quality. The quality can be checked by the following:

- Accuracy: To check whether the data entered is correct or not.
- Completeness: To check whether the data is available or not recorded.
- Consistency: To check whether the same data is kept in all the places that do or do not match.
- Timeliness: The data should be updated correctly.
- Believability: The data should be trustable.

Interpretability: The understandability of the data.

Dataset: [Aircraft Fleet Dataset](#):

## 1) Loading Data in Pandas

Step 1: Load Data in Pandas

```
import pandas as pd
df = pd.read_csv("Fleet Data.csv")
df.head()
```

|   | Parent Airline  | Airline         | Aircraft Type | Current | Future | Historic | Total | Orders | Unit Cost | Total Cost (Current) | Average Age |
|---|-----------------|-----------------|---------------|---------|--------|----------|-------|--------|-----------|----------------------|-------------|
| 0 | Aegean Airlines | Aegean Airlines | Airbus A319   | 1.0     | NaN    | 3.0      | 4.0   | NaN    | \$90      | \$90                 | 11.6        |
| 1 | Aegean Airlines | Olympic Air     | Airbus A319   | NaN     | NaN    | 8.0      | 8.0   | NaN    | \$90      | \$0                  | NaN         |
| 2 | Aegean Airlines | Aegean Airlines | Airbus A320   | 38.0    | NaN    | 3.0      | 41.0  | NaN    | \$98      | \$3,724              | 7.5         |
| 3 | Aegean Airlines | Olympic Air     | Airbus A320   | NaN     | NaN    | 9.0      | 9.0   | NaN    | \$98      | \$0                  | NaN         |
| 4 | Aegean Airlines | Aegean Airlines | Airbus A321   | 8.0     | NaN    | NaN      | 8.0   | NaN    | \$115     | \$919                | 10.3        |

## 2) Description of the dataset.

Dataset Shape: (1583, 11)  
 <class 'pandas.core.frame.DataFrame'>  
 RangeIndex: 1583 entries, 0 to 1582  
 Data columns (total 11 columns):

| #  | Column               | Non-Null Count | Dtype   |
|----|----------------------|----------------|---------|
| 0  | Parent Airline       | 1583 non-null  | object  |
| 1  | Airline              | 1583 non-null  | object  |
| 2  | Aircraft Type        | 1583 non-null  | object  |
| 3  | Current              | 859 non-null   | float64 |
| 4  | Future               | 188 non-null   | float64 |
| 5  | Historic             | 1113 non-null  | float64 |
| 6  | Total                | 1484 non-null  | float64 |
| 7  | Orders               | 348 non-null   | float64 |
| 8  | Unit Cost            | 1548 non-null  | object  |
| 9  | Total Cost (Current) | 1556 non-null  | object  |
| 10 | Average Age          | 820 non-null   | float64 |

dtypes: float64(6), object(5)  
 memory usage: 136.2+ KB

Dataset Description:

|       | Current    | Future     | Historic    | Total       | Orders     |
|-------|------------|------------|-------------|-------------|------------|
| count | 859.000000 | 188.000000 | 1113.000000 | 1484.000000 | 348.000000 |
| mean  | 24.033760  | 3.382979   | 14.513028   | 24.955526   | 26.419540  |
| std   | 41.091234  | 4.656331   | 23.763373   | 46.651526   | 43.024179  |
| min   | 1.000000   | 1.000000   | 1.000000    | 1.000000    | 1.000000   |
| 25%   | 5.000000   | 1.000000   | 3.000000    | 4.000000    | 5.000000   |
| 50%   | 12.000000  | 2.000000   | 7.000000    | 11.000000   | 13.500000  |
| 75%   | 26.500000  | 4.000000   | 16.000000   | 27.000000   | 28.250000  |
| max   | 718.000000 | 38.000000  | 325.000000  | 952.000000  | 400.000000 |

Average Age

|       |            |
|-------|------------|
| count | 820.000000 |
| mean  | 10.115000  |
| std   | 6.859362   |
| min   | 0.100000   |
| 25%   | 5.000000   |
| 50%   | 8.900000   |
| 75%   | 14.500000  |
| max   | 39.000000  |

`df.info()`: Provides an overview of the dataset, including:

- Number of rows and columns.

- Data types of each column (e.g., int, float, object).
- Number of non-null (non-missing) values in each column.

`df.describe()`: Generates summary statistics for numeric columns, such as:

- count: Number of non-missing values.
- mean: Average value.
- std: Standard deviation.
- min, 25%, 50% (median), 75%, and max: Percentile values.

**3) Drop columns that aren't useful:** Columns like Orders and Average Age which had very little data may not contribute to analysis and we can't replace them with mean as it would make a lot of data irrelevant. Removing irrelevant columns reduces complexity.

```
# Drop columns that are not useful for analysis
df = df.drop(columns=['Orders', 'Future', 'Average Age'])

# Display the first few rows after dropping columns
df.head()
```

|   | Parent Airline  | Airline         | Aircraft Type | Current | Historic | Total | Unit Cost | Total Cost (Current) |
|---|-----------------|-----------------|---------------|---------|----------|-------|-----------|----------------------|
| 0 | Aegean Airlines | Aegean Airlines | Airbus A319   | 1.0     | 3.0      | 4.0   | \$90      | \$90                 |
| 1 | Aegean Airlines | Olympic Air     | Airbus A319   | NaN     | 8.0      | 8.0   | \$90      | \$0                  |
| 2 | Aegean Airlines | Aegean Airlines | Airbus A320   | 38.0    | 3.0      | 41.0  | \$98      | \$3,724              |
| 3 | Aegean Airlines | Olympic Air     | Airbus A320   | NaN     | 9.0      | 9.0   | \$98      | \$0                  |
| 4 | Aegean Airlines | Aegean Airlines | Airbus A321   | 8.0     | NaN      | 8.0   | \$115     | \$919                |

**4) Drop rows with maximum missing values.**

```
df = df.dropna(thresh=len(df.columns) - 2)
```

- Drops rows where more than 2 columns have missing (NaN) values.

```

print(f"Rows before dropping: {df.shape[0]}")

# Drop rows with maximum missing values
df = df.dropna(thresh=len(df.columns) - 2)

# Number of rows after dropping
print(f"Rows after dropping: {df.shape[0]}")

# Display the first few rows after dropping rows
df.head()

```

Rows before dropping: 1583  
Rows after dropping: 1492

|   | Parent Airline  | Airline         | Aircraft Type | Current | Historic | Total | Unit Cost | Total Cost (Current) |
|---|-----------------|-----------------|---------------|---------|----------|-------|-----------|----------------------|
| 0 | Aegean Airlines | Aegean Airlines | Airbus A319   | 1.0     | 3.0      | 4.0   | \$90      | \$90                 |
| 1 | Aegean Airlines | Olympic Air     | Airbus A319   | NaN     | 8.0      | 8.0   | \$90      | \$0                  |
| 2 | Aegean Airlines | Aegean Airlines | Airbus A320   | 38.0    | 3.0      | 41.0  | \$98      | \$3,724              |
| 3 | Aegean Airlines | Olympic Air     | Airbus A320   | NaN     | 9.0      | 9.0   | \$98      | \$0                  |
| 4 | Aegean Airlines | Aegean Airlines | Airbus A321   | 8.0     | NaN      | 8.0   | \$115     | \$919                |

### 5)Take care of missing data.

`df.fillna(df.mean())`: Replaces missing values (NaN) in numeric columns with the mean of that column.

```

# Count missing values before filling
missing_before = df.isna().sum().sum()
print(f"Missing values before: {missing_before}")

# Fill missing values with the mean for numerical columns
# df = df.fillna(df.mean())
df = df.fillna(df.select_dtypes(include=['number']).mean())

# Count missing values after filling
missing_after = df.isna().sum().sum()
print(f"Missing values after: {missing_after}")

```

Missing values before: 1081  
Missing values after: 17

### 6)Create dummy variables.

`pd.get_dummies()`: Converts categorical variables into dummy variables (binary indicators: 0 or 1).

- Example: The Gender column becomes Gender\_Male (1 if Male, 0 otherwise).

`columns=['...']`: Specifies which columns to convert.

`drop_first=True`: Avoids the "dummy variable trap" by dropping one dummy variable to prevent multicollinearity.

|   | Parent<br>Airline  | Airline            | Current   | Historic  | Total | Unit<br>Cost | Total<br>Cost<br>(Current) | Aircraft<br>Type_ATR<br>42-<br>300F/-320F | Aircraft<br>Type_ATR<br>42-600 | Aircraft<br>Type_ATR<br>42/72 | ... | Aircraft<br>Type_McDonnell<br>Douglas MD-90 | Aircraft<br>Type_Saab<br>2000 | Aircraft<br>Type_Saab<br>340 | Aircraft<br>Type_Suk<br>Super |
|---|--------------------|--------------------|-----------|-----------|-------|--------------|----------------------------|---|--------------------------------|-------------------------------|-----|---|-------------------------------|------------------------------|-------------------------------|
| 0 | Aegean<br>Airlines | Aegean<br>Airlines | 1.000000  | 3.000000  | 4.0   | \$90         | \$90                       | False                                     | False                          | False                         | ... | False                                       | False                         | False                        | F                             |
| 1 | Aegean<br>Airlines | Olympic<br>Air     | 24.270907 | 8.000000  | 8.0   | \$90         | \$0                        | False                                     | False                          | False                         | ... | False                                       | False                         | False                        | F                             |
| 2 | Aegean<br>Airlines | Aegean<br>Airlines | 38.000000 | 3.000000  | 41.0  | \$98         | \$3,724                    | False                                     | False                          | False                         | ... | False                                       | False                         | False                        | F                             |
| 3 | Aegean<br>Airlines | Olympic<br>Air     | 24.270907 | 9.000000  | 9.0   | \$98         | \$0                        | False                                     | False                          | False                         | ... | False                                       | False                         | False                        | F                             |
| 4 | Aegean<br>Airlines | Aegean<br>Airlines | 8.000000  | 14.629091 | 8.0   | \$115        | \$919                      | False                                     | False                          | False                         | ... | False                                       | False                         | False                        | F                             |

## 7)Find out outliers (manually)

We first identified the outliers using Python by calculating the interquartile range (IQR) and marking values outside the IQR as outliers. Then, we imported the data into Excel, where we marked the outliers with a 'Yes' label. Finally, we used Excel's filtering and deletion tools to manually remove the outliers from the dataset.

```
import numpy as np

# Select only numeric columns
numeric_cols = df.select_dtypes(include=[np.number])

# Calculate Q1 (25th percentile) and Q3 (75th percentile)
Q1 = numeric_cols.quantile(0.25)
Q3 = numeric_cols.quantile(0.75)
IQR = Q3 - Q1

# Define outlier boundaries
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Find outliers
outliers = (numeric_cols < lower_bound) | (numeric_cols > upper_bound)

# Print number of outliers per column
print(outliers.sum())

print(lower_bound, upper_bound)
```

```
Current      123
Historic     131
Total        150
dtype: int64
Current     -11.406360
Historic    -11.943636
Total       -28.000000
dtype: float64 Current      45.677267
Historic      30.572727
Total         60.000000
dtype: float64
```

```
# Mark outliers: If any of the columns in the 'outliers' DataFrame are True, mark as 'Yes'
df['Outlier'] = outliers.any(axis=1).map({True: 'Yes', False: 'No'})

# Export the DataFrame with the 'Outlier' column to an Excel file
df.to_excel('marked_outliers.xlsx', index=False)

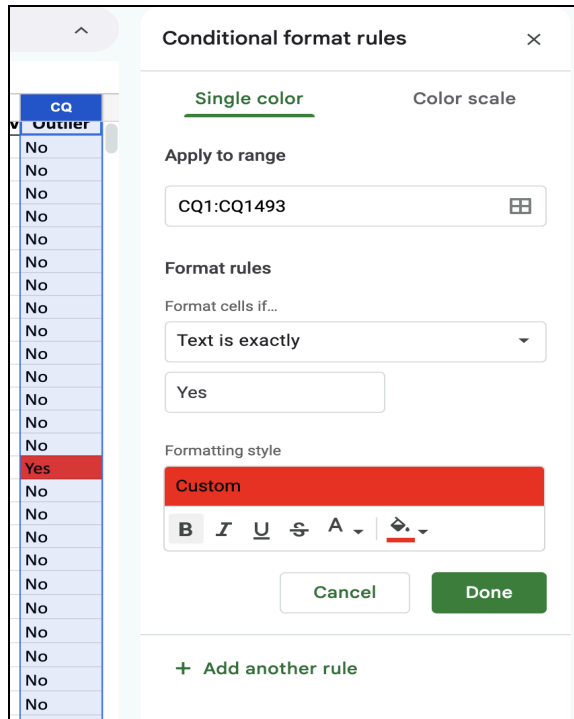
print("Outliers marked and exported to 'marked_outliers.xlsx'.")
```

```
Outliers marked and exported to 'marked_outliers.xlsx'.
```

```
# Remove rows where 'Outlier' is 'Yes'
df = df[df['Outlier'] == 'No'] # Overwrite the df to remove the outliers

# Now df is cleaned and free from outliers
print("Outliers removed, and df is updated without outliers.")
```

```
Outliers removed, and df is updated without outliers.
```



```
[15] # Remove rows where 'Outlier' is 'Yes'
df = df[df['Outlier'] == 'No'] # Overwrite the df to remove the outliers

# Now df is cleaned and free from outliers
print("Outliers removed, and df is updated without outliers.")
```

Outliers removed, and df is updated without outliers.

df.shape

(1492, 9)

## 8) standardization and normalization of columns

**Standardization** is another scaling technique where the values are centered around the mean with a unit standard deviation. This means that the mean of the attribute becomes zero and the resultant distribution has a unit standard deviation.

Standardization equation

$$X' = \frac{X - \mu}{\sigma}$$

To standardize your data, we need to import the StandardScaler from the sklearn library and apply it to our dataset.

**Normalization** is a scaling technique in which values are shifted and rescaled so that they end up ranging between 0 and 1. It is also known as Min-Max scaling.

Normalization equation

$$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$$

Here,  $X_{max}$  and  $X_{min}$  are the maximum and the minimum values of the feature respectively.

- When the value of  $X$  is the minimum value in the column, the numerator will be 0, and hence  $X'$  is 0
- On the other hand, when the value of  $X$  is the maximum value in the column, the numerator is equal to the denominator and thus the value of  $X'$  is 1
- If the value of  $X$  is between the minimum and the maximum value, then the value of  $X'$  is between 0 and 1

To normalize your data, you need to import the `MinMaxScaler` from the `sklearn` library and apply it to our dataset.

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler

# Standardization
scaler = StandardScaler()
df_standardized = pd.DataFrame(scaler.fit_transform(df.select_dtypes(include=['float64', 'int64'])), columns=df.select_dtypes(include=['float64', 'int64']).columns)

# Normalization
min_max_scaler = MinMaxScaler()
df_normalized = pd.DataFrame(min_max_scaler.fit_transform(df.select_dtypes(include=['float64', 'int64'])), columns=df.select_dtypes(include=['float64', 'int64']).columns)

# Display the first few rows after standardization and normalization
print("Standardized Data:")
print(df_standardized.head())

print("Normalized Data:")
print(df_normalized.head())
```

Standardized Data:

|   | Current   | Future        | Total     |
|---|-----------|---------------|-----------|
| 0 | -0.745112 | -5.431067e-16 | -0.454163 |
| 1 | 0.000000  | -5.431067e-16 | -0.367472 |
| 2 | 0.451790  | -5.431067e-16 | 0.347727  |
| 3 | 0.000000  | -5.431067e-16 | -0.345799 |
| 4 | -0.518671 | -5.431067e-16 | -0.367472 |

Normalized Data:

|   | Current  | Future   | Total    |
|---|----------|----------|----------|
| 0 | 0.000000 | 0.064405 | 0.003155 |
| 1 | 0.032125 | 0.064405 | 0.007361 |
| 2 | 0.051604 | 0.064405 | 0.042061 |
| 3 | 0.032125 | 0.064405 | 0.008412 |
| 4 | 0.009763 | 0.064405 | 0.007361 |



**Conclusion:**

Thus we have understood how to perform data preprocessing which can further be taken into exploratory data analysis and further in the Model preparation sequence.